Ryan Othniel Kearns

19 December 2023

Word Count: 4,999

The Extended Mind and the Integrative Stance

Introduction

There is something distinctive about my smartphone among the items in my pockets. Between the phone and its equivalent value in cash, I would much rather a pickpocket go for the cash. Here I am not appealing to the phone's sentimental value, though that explains part of my preference. My favorite family photo, stolen from my pocket, incurs sentimental suffering. Yet a stolen phone, unlike the cash or the photo, renders me *helpless*. My contact with loved ones is severed. My contactless credit cards are unavailable. My real-time translation service, interactive maps, and universal encyclopedia are offline. I am a greatly diminished version of myself without the phone. Losing it is like losing a sense. I therefore have a moral intuition that the theft of one's phone is a crueler act than the theft of equivalent cash – more like the theft of a walking cane or spectacles. These objects feel more important, more like a part of me than other everyday possessions.

Formalizing this thought, one can posit a kind of moral equivalence between certain non-bodily possessions and body parts, in that encroaching upon my relationship to these non-bodily objects constitutes an encroachment upon myself. How might we justify such a normative claim?

One attempt, from David Chalmers, employs a popular theory in the philosophy of mind: the Extended Mind Hypothesis (EMH). The EMH claims that external objects can play mental roles in governing a cognizer's perception and action. If me and my phone are causally entwined

in the right way, then my phone plays a genuinely mental role. From here, moral equivalence is easy to justify. Stealing my phone is stealing a part of my mind, which is clearly more morally severe than stealing cash.

The issue with Chalmers' approach is that the moral distinction supervenes upon a contestable metaphysical commitment. We must accept that mental states, like beliefs, extend into our environments. We must be prepared to grant mental properties to my smartphone, or my pocket calculator, or my notebook. This commitment asks a lot, and the boundaries are hard to pin down. Cases at the fringes, like pocket calculators, feel unintuitive. Moreover, the ascription of mentality to my smartphone seems like metaphysical overkill for my initial intuition, which just rested on my helplessness without the phone. Is there a path towards moral equivalence that avoids such opinionated philosophy of mind?

This paper advocates an alternative conception. I base my approach in C. Thi Nguyen's integrative stance – in short, the mental stance we adopt towards objects integrated into our goal-seeking behavior. I argue that the integrative stance justifies the aforementioned moral equivalence for extended mind objects whilst evading the metaphysical difficulties of the Extended Mind Hypothesis. In this way, we can support our normative intuition for exactly those objects that Chalmers has in mind, without unpalatable commitments to externalized mentality. The argument consists of three parts: first, that the Extended Mind Hypothesis makes objectionable metaphysical commitments; second, that the integrative stance applies to the same class of objects as the Extended Mind Hypothesis; and third, that the integrative stance unproblematically justifies the type of moral equivalence mentioned above.

Section 1 of this paper introduces the Extended Mind Hypothesis. I then raise an objection to the hypothesis in section 2, and revisit Chalmers' proposed normative implications

in section 3. Turning then to establish the integrative stance, section 4 introduces the general concept of mental stances from Peter Strawson. Section 5 focuses on C. Thi Nguyen's particular integrative stance. Section 6 ties this discussion back to the EMH, identifying a feature of extended mind objects that connects them to our discussion of mental stances. I show how the integrative stance applies to the same class of objects as the EMH in section 7. Finally, in section 8, we revisit the normative implications for these objects, and show how the ascription of mentality is not necessary for our moral equivalence claim to stand. To conclude, I discuss the implications of this argument for the philosophy of technology.

1. The Extended Mind Hypothesis

The Extended Mind Hypothesis, from Andy Clark and David Chalmers, claims that cognitive processes can extend into our environments. In order to understand which cognitive processes and which parts of our environments apply, we adopt some of Clark and Chalmers' terminology. A *coupled system* is an instance of "the human organism… linked with an external entity in a two-way interaction."[1] Consider using your pocket calculator during mathematics class. When you turn to your calculator to perform arithmetic, you input numbers into the interface, the calculator produces a result, and you act on that result. You and the calculator form a coupled system. The EMH says that coupled systems linked in this particular way are *cognitive systems*, that is, the process going on between you and your calculator is a cognitive process.

What "particular way" are Clark and Chalmers talking about? In such systems "all the components in the system play an active causal role, and they jointly govern behaviour in the same sort of way that cognition usually does."[2] In our example, your calculator plays an active role in solving the arithmetic, and its answer governs your behavior. Specifically the calculator

[1] Clark, Andy, and David Chalmers. "The Extended Mind." 8.
[2] Ibid.

plays a role usually reserved for what we might call *standard cognition* – it replaces your need for mental arithmetic. There is a functional similarity relation between this extended coupled system and an unquestionably cognitive one. Clark and Chalmers appeal to this relation as justification that said system is cognitive.

The authors' most thorough example of extended mentality is Otto. Otto is an Alzheimer's patient who has adopted a notebook to function in place of his memory. Inga is his neurotypical friend. According to the EMH, Otto's consultation of his notebook constitutes belief just as Inga's consultation of her memory does. The relevant explanatory details of each situation, and the causal roles played by the notebook and biological memory, are identical. Opponents of this view are challenged to find a relevant difference between Otto's case and Inga's, one that makes the ascription of belief in Otto's case unpalatable. It is not enough to point out that only in Inga's case is the information inside her head; this begs the question. Again, Clark and Chalmers justify their stance via the functional similarity relation between Otto's notebook and Inga's (undeniably mental) biological memory.

A well-known objection to the hypothesis concerns *reliability* or *portability*. Inga's memory is more reliably available than Otto's notebook. She cannot leave it in a taxi, and she has access to it in the shower. But note that as Otto moves throughout the world, his notebook accompanies him everywhere, and is used to achieve exactly the same goals as Inga achieves with her memory. Moreover, Inga's memory may be unavailable while asleep or intoxicated, but we do not reject her standing belief in these cases. Were Otto's notebook often missing, or sparingly employed, we should be reluctant to ascribe the EMH to his situation.

Clark and Chalmers say that Otto's case counts as one of extended belief by satisfying four essential criteria. The first is *constancy*: Otto's notebook is a constant accompaniment in his

life. Second is *availability*: the notebook is generally directly available to Otto without difficulty. Third is *endorsement*: Otto automatically endorses any information he reads from the notebook. Fourth is *past-endorsement*: Otto has consciously endorsed information in the notebook at some previous time, and this past endorsement explains the information's presence in the notebook. This fourth criterion is notably the weakest, even according to the authors. It is possible that genuine beliefs can be gained through unconscious perception or memory tampering.

Now, an objector at this point might put their foot down and say that it's just what they mean by belief that it takes place in the head. Here, Clark and Chalmers appeal to something like conceptual engineering.[3] Irrespective of the standard usage of the term, their point is that the concept of belief ought to accommodate cases like Otto's. In fact, they go as far as to say that their expanded notion of "belief" picks out "something more akin to a natural kind."[4] This statement takes a strong metaphysical stance – in my view, problematically strong, and potentially jeopardizing.

2. The Problem of Vagueness in the Extended Mind Hypothesis

I will argue that Clark and Chalmers' extended belief concept has not picked out a natural kind, due to several problems. First, it is widely rejected that consciousness could be an extended mental state. Desires and emotions might be extended, though Clark and Chalmers make only weak gestures in this direction. So mentality *eo ipso* does not demarcate this purported natural kind. Second, and more critically, the claim about a natural kind seems suspicious as many problem cases pile up at the boundaries. Several of the aforementioned essential criteria can apply in degrees. Take constancy. Otto's notebook is always with him, but what about my pocket dictionary? If my usage of the dictionary is occasional, it should not qualify. But as I move

[3] See Capellen, Herman and Plunkett, David, "A Guided Tour of Conceptual Engineering and Conceptual Ethics."
[4] Clark, Andy, and David Chalmers. "The Extended Mind." 16.

towards engaging with the book constantly in day-to-day life, and gradually inform more of my decisions with its reference, where do we draw the line? Also, my pocket dictionary initially violates the fourth criterion – it is precisely because I *don't* have the information that I am consulting it. But what if I start annotating? It becomes clear that Otto's case was a softball for purposes of explication, given how decisively it satisfies each of the four criteria. The instant we move beyond it, we see that the real work of demarcating this notion of belief is largely undone. If we stick to rigid interpretations of the four criteria, we risk absolutely nothing new falling under our concept, save fanciful thought experiments like Otto's. Alternatively, we might say that extended beliefs are sufficiently constant, sufficiently available, always endorsed, and sometimes past-endorsed, for some definitions of sufficiently and sometimes. This relaxation captures my pocket dictionary, though the vagueness of application makes predication difficult, which is precisely what we cannot have with a natural kind. We are not carving nature at the joints. Significant enough differences in degree can bring us to coupled systems that are unquestionably not belief states, such as my "believing" who won the football game before checking the news.

Beyond Otto, the case for extended beliefs appears to be on a slippery slope. In the following section, though, I acknowledge one consequence that we should desire to preserve.

3. Normative Consequences of the Extended Mind Hypothesis

The EMH may not be so easily swallowed. Yet, whether or not one accepts the hypothesis, it should be clear that the types of coupled systems picked out in this analysis are normatively distinctive. Even if Otto's notebook is not part of his mind, it is clearly disastrous to take it from him. Conversely, it is a smaller problem if my pocket dictionary is stolen. We might want to justify judging the theft of Otto's notebook more severely. "The Extended Mind" paper

devotes only a brief closing paragraph to such questions, though Chalmers in his book *Reality+* says more:

> The extended mind may reconfigure how we think about morality and the self. If someone steals my smartphone, we typically think of this as theft. But if the extended-mind hypothesis is right, it's more akin to assault. If the phone is part of me, then interfering with the phone is interfering with my person.[5]

I find this observation astute; however, are extended mental states really required for it to hold? Objects can be "part of my person" without comprising part of my mind. It would be simplifying to express this normative judgment without the evident baggage of the EMH, or even more basically, without necessitating a functionalist view of the mind. In the following sections I will show that such an expression is possible. To begin, we pivot to the philosophy of trust to introduce some new concepts.

4. Peter Strawson's Objective and Participant Stances

In "Freedom and Resentment," Peter Strawson distinguishes two kinds of *mental stances* – the objective and participant stances.[6] Strawson's machinery provides a useful framework for extended mind objects.[7] Mental stances concern the explanations we give for other moral agents' actions in the world. On the objective stance, one explains people as objects interacting with the world – or, as "causally manipulated and manipulable in various ways."[8] On the participant stance, by contrast, one regards people as fellow moral agents. Now, agents are "appropriate

[5] Chalmers, David John. "Does Augmented Reality Extend the Mind?" 307.
[6] Strawson doesn't directly use "mental stance" to refer to these concepts, though I take after Richard Holton, C. Thi Nguyen, and the Stanford Encyclopedia of Philosophy in adopting this terminology.
[7] I use "extended mind objects" as shorthand for the objects in coupled systems falling under the domain of the EMH: Otto's notebook and so on. Even if we reject the EMH, I consider the class of such objects well-formed enough to work with.
[8] Snowdon, Paul and Anil Gomes, "Peter Frederick Strawson."

objects of 'reactive attitudes.'"[9] We may direct gratitude, envy, anger, resentment, or sympathy towards such agents. Under the objective stance we can do no such thing, since the objective explanations for others' behavior, *viz* their causal manipulations of the world, do not engender such reactive attitudes. Going about the world, an agent may switch between objective and participant explanations, much as a scientist cycles through microscope lenses. As these two mental stances often involve the same subject matter – the actions of other moral agents – we may have both objective and participant explanations for the same action. For an analogous case in the philosophy of science, consider Wilfrid Sellars' manifest and scientific images.[10]

Richard Holton applies Strawson's participant stance to motivate reactive attitudes that hold distinctively in cases of trust between individuals.[11] Following Annette Baier in the philosophy of trust,[12] Holton is interested in the characteristics separating trust from reliance. Holton says that trust involves reliance plus a particular reactive attitude: the readiness to feel betrayal should one's trust be broken. Here, Strawson's mental stance machinery helps distinguish trust from reliance. One may take the objective stance towards another they rely on, since objective explanations alone suffice to characterize such reliance. However, when I trust you, I necessarily take the participant stance towards you. I see you as a fellow moral agent, and so I prepare myself with the appropriate reactive attitude – the willingness to feel betrayed if my trust is broken.

For an illustrative example of this distinction, consider two similar cases. Suppose you and I have neighboring offices, and both teach classes at the same time each morning. In the first case, I have taken your leaving for class as a cue to leave for my own class. Whenever you

---

[9] Ibid.
[10] Sellars, Wilfrid. "Philosophy and the Scientific Image of Man."
[11] Holton, Richard. "Deciding to Trust, Coming to Believe."
[12] Baier, Annette. "Trust and Antitrust."

happen to pass by my open office door in the morning, I pack up and leave so as to also arrive on time. I *rely* on your passing by my door to get me to class on time. I can explain your actions under the objective stance, since all I rely on is the physical fact of your passing by my door. In the second case, however, I have asked you to knock on my door to remind me to leave. Now, I *trust* you to remind me to leave on time. In this second case your behavior is explainable under the participant stance, since I must take into account your own moral agency, your intentions, and the conditions of our cooperative relationship. Here is where betrayal comes in. Suppose that one morning you fall ill and stay home for the day, and that I am not made aware of your absence. In the first case, where you simply fail to pass by my open door, I surely cannot blame you for my own tardiness to class. I cannot feel *betrayed* by your failure to pass by the door. In the second case, though, I can feel betrayed. Under the participant stance, I have particular normative expectations about your reminding me to get to class on time. These expectations make possible the reaction of betrayal.

5. C. Thi Nguyen's Integrative Stance

Like Richard Holton, C. Thi Nguyen seeks to disambiguate trust and reliance. He similarly roots his analysis in the differing mental stances afforded to trust and reliance. However, Nguyen motivates a different mental stance as the basis for trust. Nguyen thinks that trust is a kind of *unquestioning attitude* we take towards objects in our environment. We trust precisely when we refrain from questioning the reliability of objects. In familiar cases these objects will be other human actors. Trusting your pharmacist means not considering whether they have laced your prescription with cyanide. Trust can permit us to take actions we would not otherwise – to take a pill, or to disclose sensitive medical history.

The unquestioning attitude also applies to objects. A climber trusting their rope puts questions of the rope's integrity out of mind, and leaps confidently to the next hold. When we follow the directions of our mapping software without question, we are trusting the software. When we reactively step on our brakes we are trusting our car. In Nguyen's view, we place trust in certain objects exactly as we do in people.

To some, this view prompts immediate scrutiny. Recall that the goal is to distinguish trust from reliance. Has Nguyen succeeded in this task? Some objectors say no. Certainly we *rely* on our car's brakes, but we don't *trust* them! After all, how does it make sense to say that things like our climbing ropes and automobiles can betray us?

To face this charge, Nguyen advances a new kind of mental stance: the *integrative stance*. Under the integrative stance, we consider teleological explanations for the roles of objects in our environment. The objective stance presents your car brakes as a system of disks and ablative plating, but a teleological explanation highlights the goal-directed aspect of this system: its function to stop your car. As agency is goal-directed, teleological explanations can articulate how we integrate the actions of objects into our own agency.

Nguyen argues that this integration can bring about the reactive attitude of betrayal. The argument goes thus: agents in the world have particular goals. Achieving certain goals requires reliance on other agents or objects for actions we cannot take ourselves. This observation about modern life is uncontroversial. Filing taxes requires attorneys; eating healthy requires farmers; navigating new cities requires a smartphone. We might seek to supervise these agents and objects to safeguard our reliance on them. Yet in most practical cases, complete supervision is impossible – auditing a tax attorney's full-time job is, itself, a full-time job. Where we fail to supervise, we may instead aim to trust, which involves adopting an unquestioning attitude

towards the actions and decrees of our attorneys, farmers, and smartphones. In these cases, we trust in order to incorporate those actions and decrees into our own agency. We relegate steps toward our goals to the function of external objects. Such relegation incurs vulnerability, since our goals matter to us, and we must cede some control over them. This vulnerability primes us with the right sorts of reactive attitudes for betrayal. We bring external objects into our agency, and when they fail us, curtailing our agency, we feel betrayed.

Nguyen says that the "normative bite" that attends such situations stems from our desire to integrate external objects into our agency.[13] As to whether the particular normative attitude Nguyen has identified is *betrayal*, I have doubts.[14] Yet I think the weaker claim that *some normative distinction is afoot here* deserves to be taken seriously. When my car brakes suddenly fail, I do not just explain the situation under Strawson's objective stance. That is, I do not only regard my car brakes as objects "causally manipulated and manipulable in various ways." Instead, I have a teleological explanation in mind. The function of my car brakes is tightly integrated into my agency and my goals, such as getting to work on time and not dying in a car crash. I switch lenses, and consider my car brakes as objects deserving of reactive attitudes, much like other, more cognitive, objects at my disposal, such as my memory. What unifies my attitudes towards these objects is not their mentality, but the extent of my functional integration with them. In the next section I will make this idea of functional integration precise.

---

[13] Nguyen, C. Thi. "Trust as an unquestioning attitude." 30.

[14] I do not think that Nguyen's unquestioning attitude is sufficient to characterize trust. I advance an objection in "Contextual Trust," my undergraduate thesis, where I show that Nguyen's account ascribes trust in unintuitive situations, such as trust directed towards individuals one despises. That full objection is out of scope and unrelated to my use of the integrative stance here.

If I do not accept that the unquestioning attitude characterizes trust, it follows that I need not commit to the possibility of betrayal in cases where the unquestioning attitude applies. Here I grant the widely accepted claim in the philosophy of trust that trust brings along the possibility of betrayal, as stated by Annette Baier in her 1986 essay, "Trust and Antitrust."

6. Two Kinds of "Coupled Systems"

Returning to the EMH, I contend that the integrative stance helps pick out the normative distinction afforded to extended mind objects. Before advancing this argument, though, I will make a distinction about coupled systems that is missing in Clark and Chalmers' initial treatment of the subject.

Recall that Clark and Chalmers motivate the idea of *reliable coupling* to defend their view from the portability objection. The objection seeks to disqualify extended mind objects because the coupled systems they form are easily decoupled. How can a notebook count as a part of my working memory if I could so easily leave it in a taxicab? True coupled systems in the extended mind, so argue Clark and Chalmers, are not coupled in any flippant way. Instead, coupled properly, their reliability starts to approach that of the biological brain. This quality invokes the first two essential criteria of extended beliefs, constancy and availability. Brain parts happen to be both constant and available, and hence reliably coupled, but not only brain parts can be reliably coupled.

Reliability, however, only gets us partway to specifying the kinds of coupled systems that Clark and Chalmers care about. They neglect to mention an important orthogonal characteristic to reliability, which I will call *dependence*. Note that reliability has to do with the physical or structural integrity and robustness of the coupling. Brain parts are typically reliably coupled, and notebooks are typically not. A second, orthogonal property, dependence, measures the extent of functional integration for a given coupled system. Coupled systems can have high or low reliability and, separately, high or low dependence. The human brain is an instance of both reliable and dependent coupling, while one's car brakes exemplify just dependent coupling, and the human appendix exemplifies just reliable coupling – always with you, always available, but

never for anything functionally useful.[15] Note that dependence implicates the third and fourth criteria for extended mind objects, endorsement and past-endorsement. Otto's notebook is only populated with information he consciously endorses (past-endorsement), and he automatically accepts old information in the notebook (endorsement). These criteria explain the high functional integration of this coupled system. The reliability of the notebook, meanwhile, is orthogonal. For another example illustrating this orthogonality consider networked computer systems, where programs are directly dependent upon one another yet can be both disconnected or made unavailable due to network issues.

While Clark and Chalmers defend the importance of reliable coupling for the extended mind, they do not discuss dependent coupling, and this is an omission. Every example in "The Extended Mind" describes a goal-oriented functional process. Via cognitive extension into the environment, the mind can rotate blocks in Tetris, or find words in Scrabble tiles, or perform long division. Clearly, such systems will have varying levels of dependent coupling. This variance, again, can be independent of the reliability of the coupling. Why does this matter? Bringing back the integrative stance, I will argue that our mental stance towards extended mind objects depends on how dependently (and *not* reliably) they are coupled.

7. A Convergence Thesis for the Extended Mind Hypothesis and the Integrative Stance

It is now possible to advance an argument connecting the domain of objects covered in the Extended Mind Hypothesis with those explainable under the integrative stance. In section 5 we proposed the integrative stance, a mental framework that ends-directed agents adopt towards objects they employ in the pursuit of their goals. In section 6, we saw two orthogonal dimensions for qualifying the integration of such objects in coupled systems. Reliability specifies the

[15] I hope that recent research into the immune function of the appendix can be ignored in the spirit of this point. We can devise other examples of just reliable coupling, like purely cosmetic tattoos.

difficulty of physically separating the integrated object, and dependence specifies the extent of functional integration.

Dependence, not reliability, motivates the integrative stance. Since the integrative stance has to do with functional integration, and functional integration is independent of physical reliability of coupling, it follows that reliability should not affect the integrative stance. The integrative stance does not admit explanations having to do with the reliability of the integration. One might be foolish to depend on an object with low reliability – navigating with glitchy mapping software is foolish – but this practical qualm does not change *what it is* to adopt the stance.

So, the integrative stance applies to cases where one has high dependence on an integrated object. Since the stance highlights an object's teleological integration, and dependence measures the teleological or functional integration of an object, the integrative stance will offer a suitable explanatory framing for such objects.

Separately, I argue that if the Extended Mind Hypothesis applies to an integrated object, then one has high dependence on that object. In section 6 I claimed that high dependence implicates the endorsement and past-endorsement criteria. This claim is slightly stronger, saying that the four criteria in the EMH imply high dependence. While I have shown that Clark and Chalmers' four criteria are problematic,[16] I take this to be an intention of the hypothesis. Indeed, Clark and Chalmers do not raise a single example for the extended mind where dependence seems to be low. By contraposition, it seems likely that a coupled system exhibiting low dependence would not count as a case of extended mentality. The hypothesis is a functionalist

[16] My argument does not require the EMH to go through, provided the domain of the hypothesis is reasonably well defined. Note that unlike Clark and Chalmers in the case of extended belief, I am not proposing that "coupled systems exhibiting dependence" identifies a natural kind. Dependence can vary in degree. Thus, interpretations of the EMH at varying levels of rigidity should all work for my argument.

one, and what mentality is there in a coupled system where the parts play no integrated functional role?

If we accept that the EMH identifies coupled systems exhibiting high dependence, and that the integrative stance captures cases of high dependence in coupled systems, it follows that the domain of the EMH converges with the integrative stance. There is explanatory virtue in the integrative stance adopted towards objects in the extended mind.

8. Normative Distinctions, Revisited

What do we gain from the conclusion that the domain of the EMH and the integrative stance converge? Recall Chalmers' attempt at a normative result for extended mind objects, covered in section 3. One of the peripheral objectives of the EMH was to "reconfigure how we think about morality and the self."[17] The integrative stance has accomplished this aim for the same class of objects that Chalmers considers. It explains and justifies my initial intuition. Returning to our opening thought experiment: suppose I am new to New York City, and using my smartphone to navigate around. I am accustomed to immediately endorsing the navigation instructions of Google Maps, and critically I have no ability to navigate in this way using my brain. I am an agent with goals – to navigate to the museum, say, and to not get lost – and under the integrative stance, my smartphone is a tightly integrated feature of my environment that is critical to my achieving these goals. To steal my smartphone in this circumstance is to leave me both epistemically and practically vulnerable, and is a clear violation of my agency. In fact, that my smartphone might form part of my mind seems a weak normative justification by comparison. The smartphone's integration with my agency directly explains why interference with it is interference with myself as an agent. We have achieved Chalmers' objective to show

[17] Chalmers, David John. "Does Augmented Reality Extend the Mind?" 307.

how “interfering with the phone is interfering with my person,” and we need not uphold any contentious opinions about the mind to say so.

Conclusion

Our conclusion in this paper is that the integrative stance applies to extended mind objects. We should uphold normative expectations towards such objects that surpass the expectations we hold of generic objects. Like the EMH, this conclusion may support a new way of thinking about cognitive enhancers like Google search, mapping softwares, and augmented reality technologies. Rather than merely tools, we should consider such technologies as functional extensions of the self, subject to the same care and alarm that we use to attend to our own bodies and brains. In addition, this framework should be acceptable to skeptics of the Extended Mind Hypothesis, who feel that such commitments in the philosophy of mind are unwarranted.

Works Cited

Baier, Annette. “Trust and Antitrust.” *Ethics 96*, no. 2 (1986): 231–60. https://doi.org/10.1086/292745.

Chalmers, David John. “Does Augmented Reality Extend the Mind?” Essay. In *Reality+: Virtual Worlds and the Problem of Philosophy*, 294–309. London: Penguin Books, 2022.

Clark, Andy, and David Chalmers. “The Extended Mind.” *Analysis 58*, no. 1 (1998): 7–19. http://www.jstor.org/stable/3328150.

Holton, Richard. “Deciding to Trust, Coming to Believe.” *Australasian Journal of Philosophy* 72, no. 1 (1994): 63–76. https://doi.org/10.1080/00048409412345881.

Nguyen, C. Thi. “Trust as an unquestioning attitude.” *Oxford Studies in Epistemology* (2020): 1–45.

Sellars, Wilfrid. “Philosophy and the Scientific Image of Man.” *Frontiers of Science and Philosophy*, Robert Colodny (ed.) (Pittsburgh, PA: University of Pittsburgh Press, 1962): 35–78.

Snowdon, Paul and Anil Gomes. “Peter Frederick Strawson”. The Stanford Encyclopedia of Philosophy (Summer 2021 Edition), Edward N. Zalta (ed.). https://plato.stanford.edu/archives/sum2021/entries/strawson/.

Strawson, Peter Frederick. “Freedom and Resentment.” *Proceedings of the British Academy* 48 (1962): 1–25.